

# Learning-Guided Simulated Annealing for the Capacitated Vehicle Routing Problem

Anonymous submission

## Abstract

## Introduction

## Background and Related Work

### Problem Definition

The *Capacitated Vehicle Routing Problem (CVRP)* is defined on a complete undirected graph  $G = (V, E)$ , where  $V = \{0\} \cup \mathcal{C}$  and  $E = V \times V$  denote the sets of nodes and edges, respectively. Node 0 corresponds to the central depot, and  $\mathcal{C} = \{1, \dots, N\}$  represents the set of  $N$  customers.

Each edge  $(i, j) \in E$  is associated with a travel cost  $c_{ij} = \|\mathbf{p}_i - \mathbf{p}_j\|_2$ , the Euclidean distance between the 2D coordinates  $\mathbf{p}_i, \mathbf{p}_j \in [0, 1]^2$ , satisfying  $c_{ij} = c_{ji}$ . Each customer  $i \in \mathcal{C}$  has a non-negative demand  $q_i$ , while the depot has zero demand ( $q_0 = 0$ ). An unlimited fleet of identical vehicles, each with uniform capacity  $Q$ , is stationed at the depot; the number of routes  $m$  is therefore a decision variable unconstrained from above.

A route  $R_k = (0, v_1, v_2, \dots, v_p, 0)$  is a simple cycle starting and ending at the depot, where  $\{v_1, \dots, v_p\} \subseteq \mathcal{C}$  denotes the subset of customers served by vehicle  $k$ . A *solution*  $s = \{R_1, \dots, R_m\}$  consists of a set of  $m$  routes. A solution  $s$  is *feasible* if and only if it satisfies the following conditions:

1. Each customer is visited exactly once by a single vehicle,

$$\sum_{k=1}^m \mathbf{1}_{\{i \in R_k\}} = 1, \quad \forall i \in \mathcal{C}.$$

2. For every route  $R_k$ , the total demand of its customers does not exceed the vehicle capacity  $Q$ , i.e.,

$$\sum_{i \in R_k \setminus \{0\}} q_i \leq Q, \quad \forall k = 1, \dots, m.$$

The set of *feasible solutions* is denoted by  $\mathcal{F}$ .

Let  $x_{ij}(s)$  be a binary indicator variable with  $x_{ij}(s) = 1$  if edge  $(i, j)$  is traversed in any route  $R_k$  of  $s$ , and  $x_{ij}(s) = 0$

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

otherwise. The objective of the CVRP is to find a feasible solution  $s^* \in \mathcal{F}$  that minimizes the total travel cost  $C(s)$ :

$$s^* = \arg \min_{s \in \mathcal{F}} \left\{ C(s) = \sum_{i \in V} \sum_{j \in V, i < j} c_{ij} x_{ij}(s) \right\}.$$

## Proposed Method

### Overview

We propose *Learning-Guided Simulated Annealing (LG-SA)*, a hybrid method that combines *Simulated Annealing (SA)* with a lightweight neural policy trained by *Reinforcement Learning (RL)*. Following the Neural Simulated Annealing framework (Correia, Worrall, and Bondesan 2023), the SA component orchestrates the global exploration of the solution space through its acceptance criterion and cooling schedule, while the neural network replaces the single blind component of SA — the uniform move proposal — with a learned, state-dependent proposal distribution. This design contrasts with “heavyweight” improvement architectures such as the Dual-Aspect Collaborative Transformer (Ma et al. 2021), which replace the heuristic logic entirely with a deep encoder: our policy remains deliberately small, so that the guided search retains a per-step cost close to that of classical SA and can be batched over thousands of instances on a single GPU.

*Simulated Annealing* (Kirkpatrick, Gelatt, and Vecchi 1983) iteratively explores the neighborhood of a current solution. Given a solution  $s$ , SA samples a neighboring solution  $s' \in \mathcal{N}(s)$  and accepts it according to the *Metropolis criterion*: improving moves are always accepted, while worsening moves are accepted with probability  $\exp(-\Delta C/T)$ , where  $\Delta C = C(s') - C(s)$  is the cost difference and  $T$  the current temperature. The temperature decays from an initial value  $T_0$  to a final value  $T_{\text{final}}$  over the search horizon according to a monotone cooling schedule, which regulates the transition from exploration to exploitation.

Neighboring solutions are produced by a *neighborhood operator*  $H : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$  that maps a solution  $s$  and an action  $a \in \mathcal{A}(s)$  to the neighbor  $s' = H(s, a)$ , so that

$$\mathcal{N}(s) = \{H(s, a) : a \in \mathcal{A}(s)\}.$$

In our setting, an action designates a pair of nodes  $a = (i, j)$ , and  $H$  applies a standard pairwise local search move to them (e.g., relocating node  $i$  next to node  $j$ ). The operator is a

---

**Algorithm 1** Learning-Guided Simulated Annealing

---

**Input:** instance, initial solution  $s$ , policy  $\pi_\theta$ , operator  $H$ **Output:** best solution found  $s^*$ 

```
1: Initialize  $s^* \leftarrow s$  and temperature  $T \leftarrow T_0$ 
2: while termination criterion not met do
3:   Encode  $s$  into per-node feature vectors
4:   Sample action  $a \sim \pi_\theta(\cdot \mid s)$  (guided proposal)
5:   Generate candidate  $s' \leftarrow H(s, a)$ 
6:    $s \leftarrow s'$  with probability  $\min(1, e^{-\Delta C/T})$  (Metropolis)
7:   Update  $s^*$  if  $C(s) < C(s^*)$ ; cool  $T$ 
8: end while
9: return  $s^*$ 
```

---

modular component of the framework; the candidate operators and their comparison are presented in the experimental section.

Whereas classical SA draws  $a$  uniformly over  $\mathcal{A}(s)$ , LG-SA samples it from a policy  $\pi_\theta$  conditioned on the current solution:

$$a \sim \pi_\theta(\cdot \mid s) \quad \text{and} \quad s' = H(s, a) \in \mathcal{N}(s).$$

Everything else — the operator, the Metropolis acceptance, the cooling schedule — is kept from classical SA. The overall procedure is summarized in Algorithm 1 and illustrated in Figure 1.

### The Guided Search as a Markov Decision Process

To learn the proposal distribution with RL, we frame the SA loop as a Markov Decision Process. At step  $t$ , the *state*  $s_t$  is the current feasible solution together with the search meta-information (current temperature and remaining budget), encoded as a set of per-node feature vectors described below. The *action*  $a_t = (i, j)$  is the node pair handed to the operator  $H$ . The *transition* is the composition of the operator and the stochastic Metropolis acceptance: the next state is  $H(s_t, a_t)$  if the move is accepted and  $s_t$  otherwise, so the environment dynamics include the acceptance randomness of SA.

The *reward* is derived from the evolution of the solution cost along the search: a move that is accepted and improves the solution receives a positive reward proportional to the (normalized) cost reduction, a rejected move yields no improvement signal, and proposals leading to infeasible neighbors can be penalized. Reward variants that instead emphasize improvements of the best solution found so far, or the final solution quality at the end of the horizon, fit the same framework; the precise shaping is a design choice studied experimentally. The agent thus maximizes the expected cumulative discounted reward over the SA horizon, which aligns the learned proposal with the overall goal of driving the annealing process toward low-cost solutions rather than with any single greedy step.

### Node-Centric State Representation

To be applicable to any problem size  $N$  and to remain *scale-invariant*, the state is encoded in a *node-centric* fashion:

each node  $i$  carries its own local feature vector  $\mathbf{x}_i$ , and no fixed-size global descriptor is used. The features fall into four groups:

- *Static instance features*: node coordinates, polar coordinates relative to the depot, a depot indicator, and the normalized demand  $q_i/Q$ ;
- *Solution topology and local geometry*: the coordinates of the node’s current predecessor and successor in its route, together with local quality signals such as the detour cost incurred by the node at its current position, its distance to the route centroid, and local density statistics of the instance;
- *Capacity status*: the load ratio and remaining slack of the node’s current route, and the node’s own contribution to that load;
- *Search meta-features*: the current (normalized) SA temperature and the fraction of the search budget remaining, which let the policy adapt its behavior to the phase of the annealing schedule.

All features are local to a node and standardized, so the same trained model can be deployed on instances of arbitrary size without retraining. The complete list of features, with their formal definitions, is given in the appendix; their individual contribution is analyzed in the experimental section.

### Two-Stage Conditional Policy

The policy must define a distribution over the  $\mathcal{O}(N^2)$  node pairs while remaining cheap enough to be evaluated at every SA step. We therefore factorize the joint selection as

$$P(i, j \mid s) = P(i \mid s) \cdot P(j \mid i, s),$$

and implement it with a lightweight two-stage architecture, illustrated in Figure 2. In the first stage, a small network  $f_\theta$  maps each feature vector to a scalar score  $z_i = f_\theta(\mathbf{x}_i)$ , and the first node is sampled from the softmax distribution over these scores:

$$i \sim P(\cdot \mid s) = \text{softmax}(\{z_i : i \in \mathcal{C}\}).$$

In the second stage, the representation of the selected node  $i$  is combined with that of every other node  $j$ , and a second network  $g_\theta$  assigns a conditional score  $z_{ij} = g_\theta(\mathbf{x}_i, \mathbf{x}_j)$  from which the partner is sampled:

$$j \sim P(\cdot \mid i, s) = \text{softmax}(\{z_{ij} : j \in \mathcal{C}\}).$$

Both stages are built from small multilayer perceptrons applied identically to every node, which makes the policy permutation-equivariant. The second stage can additionally be conditioned on cheap pairwise descriptors of the candidate move (e.g., the exact cost of the move joining  $i$  and  $j$ ), computed only for the selected  $i$ .

This factorization also provides a natural hook for the capacity constraints: once  $i$  is selected, the partners  $j$  for which  $H(s, (i, j))$  would be infeasible are masked out of the second-stage distribution, so every proposed neighbor is feasible at the price of  $\mathcal{O}(N)$  feasibility checks per step. Alternative constraint-handling strategies are discussed and compared in the experimental section.




Figure 1: The LG-SA loop. Classical SA components (gray) are kept unchanged: the operator  $H$ , the Metropolis acceptance, and the cooling schedule. The only modified block is the move proposal (orange): the current solution is encoded into per-node feature vectors and a learned policy samples the node pair  $a_t = (i, j)$ , replacing the uniform sampling of classical SA.




Figure 2: Two-stage conditional sampling of the node pair  $(i, j)$ , shown on a toy route of three customers. Stage 1: every node’s feature vector is scored independently by a shared network  $f_\theta$  and the first node  $i$  is sampled from the softmax over the scores. Stage 2: the representation of the selected node  $i$  is combined with that of every candidate partner  $j$  and scored by a second network  $g_\theta$ ; infeasible partners are masked out before the softmax, and  $j$  is sampled from the conditional distribution. The factorization  $P(i, j) = P(i) \cdot P(j | i)$  makes each proposal  $\mathcal{O}(N)$ .

Since the conditional network is evaluated only for the single selected first node, each proposal costs  $\mathcal{O}(N)$  — as opposed to the  $\mathcal{O}(N^2)$  cost of scoring all pairs jointly — and the whole step is expressed as batched tensor operations, allowing thousands of instances to be advanced in parallel on one GPU.

## Training with PPO

The policy is trained with *Proximal Policy Optimization* (PPO) (Schulman et al. 2017), the standard on-policy actor-critic algorithm. Training alternates between a *collection* phase, in which the current policy runs the full LG-SA loop of Algorithm 1 on a large batch of freshly generated instances in parallel and stores every transition  $(s_t, a_t, r_t)$ , and an *update* phase, in which the stored episodes are replayed for several optimization epochs. Advantages are estimated with Generalized Advantage Estimation, and the actor maximizes the usual clipped surrogate objective with an entropy bonus; we refer to Schulman et al. (2017) for the standard formulation and only detail our departures from it.

The critic  $V_\phi$  is a separate permutation-invariant network: it encodes each node feature vector independently, pools the embeddings across nodes into a fixed-size summary, and maps this summary to a scalar value estimate. Keeping the

actor and critic distinct prevents interference between the policy and value representations.

Our implementation departs from standard PPO in three ways, all aimed at stabilizing training on the long and noisy SA horizon:

- *Hybrid clipping–KL objective.* Rather than choosing between the clipped surrogate and the KL-penalized objective, we use both: an adaptive Kullback–Leibler penalty  $\beta_{\text{KL}} D_{\text{KL}}(\pi_{\theta_{\text{old}}} || \pi_\theta)$  is added to the clipped loss, and its coefficient is doubled or halved after each update epoch depending on whether the measured divergence leaves a band around a target value. This soft trust region allows many optimization epochs on each collected batch without catastrophic policy drift.
- *KL-coupled learning rate.* The same divergence signal modulates the actor’s learning rate within fixed bounds — gently decreased when the policy moves too fast, increased when it stalls — so the effective step size tracks the trust region automatically.
- *Regularized critic.* The value loss is clipped, mirroring the actor’s trust region, and augmented with a gradient penalty that encourages a smooth (near-Lipschitz) value landscape; we found this useful given that consecutive

states differ by a single local move whose acceptance is itself stochastic.

## Experimental Results

### Experimental Setup

**Benchmark.** All experiments reported in this section use the uniform benchmark distribution of Nazari et al. (2018): depot and customer coordinates drawn uniformly in  $[0, 1]^2$ , demands drawn uniformly in  $\{1, \dots, 9\}$ , and a size-dependent vehicle capacity ( $Q = 20, 30, 40, 50$  for  $N = 10, 20, 50, 100$ ). For each size we pre-generate a fixed test set of 10,000 instances that is reused by every configuration, so that all comparisons are paired at the instance level. Unless stated otherwise, results are reported at  $N = 100$ .

**Evaluation protocol.** A trained model is evaluated by running the LG-SA loop over the full test set in parallel on a single GPU, starting from random feasible solutions and running 10,000 SA steps per instance with half-precision inference. We report the mean final routing cost over the 10,000 instances and the wall-clock time of the whole batch. Classical (blind) SA — uniform proposals with the same operator, step budget, and cooling schedule — is evaluated under the identical protocol, as is an OR-Tools baseline with a fixed per-instance time limit (Furnon and Perron 2024).

**Statistical protocol.** Every configuration is trained under 3 random seeds (5 for confirmation and headline runs), and we report means over seeds. The seed-to-seed spread of the reference configuration, measured once by repeated training and evaluation of the baseline, serves as the empirical noise floor: differences below this threshold are treated as ties.

**Sequential validation protocol.** The design space of LG-SA — operator, feasibility strategy, features, architecture, reward, annealing parameters, initialization, optimization budget, curriculum — is far too large for a full factorial study. We therefore validate it by coordinate descent: the axes are swept one at a time in order of expected leverage, and the winner of each sweep is frozen into the configuration before the next axis is examined, so later sweeps inherit all earlier decisions. Targeted follow-up batches probe the interactions we consider most at risk (e.g., the construction of the minimal feature set and the curriculum mechanism ablation), and the complete winning configuration is re-trained from scratch with 5 seeds as a final confirmation.

### Local Search Operators and Feasibility Strategies

The first axis fixes the foundational search mechanics: the operator  $H$  and the way capacity constraints are enforced.

**Operators.** We consider three standard pairwise operators, illustrated in Figure 3:

- *Swap*: exchanges the positions of nodes  $i$  and  $j$ , within a route or across routes;
- *Insertion (relocation)*: removes node  $i$  from its current position and reinserts it immediately after node  $j$ ;
- *2-opt / 2-opt\**: for intra-route pairs, reverses the path segment between  $i$  and  $j$ ; for inter-route pairs, exchanges the route suffixes following  $i$  and  $j$  (tail swap).

| Operator  | Masking             | Indirect     | Rejection    |
|-----------|---------------------|--------------|--------------|
| Insertion | <b>16.68 ± 0.03</b> | 16.97 ± 0.04 | 26.36 ± 0.71 |
| Swap      | 19.50 ± 1.54        | 22.82 ± 1.36 | 35.75 ± 3.04 |
| 2-opt(*)  | 18.93 ± 0.09        | 21.29 ± 0.39 | 25.84 ± 0.42 |

Table 1: Operator  $\times$  feasibility study: mean final cost  $\pm$  seed std at  $N = 100$  ( $10^4$  SA steps, 10,000 instances, 3 seeds; all runs start from mean cost 57.87). Insertion with proactive masking wins and is also the most stable cell.

The operator is fixed for a given model: the policy is trained to propose moves for one operator, and learns a distribution tailored to it.

**Feasibility strategies.** Under capacity constraints, many actions produce an infeasible neighbor  $H(s, a) \notin \mathcal{F}$  — the central obstacle in extending neural SA from unconstrained problems to the CVRP. We compare three ways of keeping the search inside  $\mathcal{F}$ :

- *Proactive filtering (action masking)*. The policy is restricted to the feasible subset  $\mathcal{A}_f(s) = \{a : H(s, a) \in \mathcal{F}\}$  by masking infeasible actions before sampling. The conditional factorization makes this cheap: the first node  $i$  is selected freely, and only then are the feasible partners  $\{j : H(s, (i, j)) \in \mathcal{F}\}$  computed and used to mask the second stage —  $\mathcal{O}(N)$  checks per step instead of  $\mathcal{O}(N^2)$ . Every proposal is feasible by construction; since the fleet is unlimited, a node can always open a new route, so the feasible set is never empty.
- *Reactive rejection*. The policy samples any action from its learned distribution and the resulting neighbor is checked afterwards: if  $s' \notin \mathcal{F}$ , the move is rejected and the current solution retained. This avoids the filtering overhead but can waste both compute and learning signal when many sampled moves are infeasible.
- *Indirect representation*. Feasibility is delegated to the representation: the solution is encoded as an unconstrained permutation of the customers, the policy modifies the permutation, and a deterministic *greedy split* heuristic packs it back into capacity-feasible routes (Vidal 2016). Every state is feasible by construction and the policy never handles constraints, but the reconstruction runs at every step and its cost grows with the problem size.

We train one policy for each of the nine operator–strategy combinations under the common protocol (3 seeds per cell, 27 runs); the winning pairing is carried unchanged through all subsequent experiments. Table 1 reports the full grid; the complete set of diagnostic plots (heatmap, cost–compute trade-off, per-seed spread) is given in the appendix.

**Results.** Insertion with proactive masking is the best cell (16.68) and *dominates*: it attains the lowest cost, the lowest seed variance (std 0.03, against up to 3.0 for the swap cells), and is among the cheapest in compute (173s for the full 10,000-instance evaluation). The two factors do not interact perversely — the best operator and the best strategy also combine into the best cell. Averaged over feasibility strategies, insertion (20.00) beats 2-opt (22.02) and swap (26.03):




Figure 3: The three pairwise operators on two route fragments (top: before, bottom: after). *Swap* exchanges  $i$  (orange) and  $j$  (green). *Insertion* removes  $i$  and reinserts it immediately after  $j$ . *2-opt\** exchanges the route suffixes following  $i$  and  $j$ ; when  $i$  and  $j$  belong to the same route, the operator instead reverses the path segment between them (standard 2-opt).

relocating a single node is the move that best matches a learned pointer to a promising node pair, whereas the value of a swap or a segment reversal depends more delicately on both endpoints simultaneously, making the credit assignment noisier — visible in the swap column’s large seed variance. Averaged over operators, masking (18.37) edges out the indirect representation (20.36), while reactive rejection is far behind (29.31) and is uniformly the worst choice for every operator: discarding infeasible proposals wastes both search budget and learning signal, and the policy must spend capacity learning the feasibility boundary instead of move quality. The indirect (greedy-split) representation is competitive in cost for insertion but pays the reconstruction at every step; 2-opt under masking is the slowest cell (301 s, about  $1.7\times$  the rest) because suffix exchanges require more expensive feasibility checks, yet still loses to insertion on quality — its extra compute buys nothing.

**Adopted configuration.** All subsequent experiments use the *insertion* operator with *proactive masking*. The decision is robust: the runner-up (insertion with the indirect representation, 16.97) points to the same operator, so even if the two feasibility strategies were statistically tied the operator choice would stand.

### State Feature Ablation

The node representation is validated with three complementary protocols, all trained on the winning operator–feasibility pairing and all operating on the feature groups defined in the appendix. *Leave-one-out* removes a single group from the full set and measures the degradation; *thematic* ablations remove a whole correlated cluster at once (e.g., all three density scales, or all capacity signals), exposing groups whose members are individually redundant but collectively useful; *isolation* adds a single group to a minimal base containing only the static and meta features, measuring its standalone value. From these measurements we construct candidate minimal feature sets — keeping groups whose removal hurts, dropping those that fail both the leave-one-out and the isolation test, and keeping one representative per correlated cluster — and confirm the candidates seed-paired against the full set. Finally, the conditional pair descriptors of the second stage (edge distance-ranks and exact insertion cost, defined in the

| Group                        | $\Delta_{\text{LOO}}$ | $\Delta_{\text{iso}}$ |
|------------------------------|-----------------------|-----------------------|
| Static                       | <b>+3.84 ± 0.85</b>   | (in base)             |
| Detour $\Delta_i$            | <b>+0.60 ± 0.08</b>   | <b>+2.60 ± 0.07</b>   |
| Topology                     | <b>+0.46 ± 0.08</b>   | <b>+2.19 ± 0.15</b>   |
| Meta                         | <b>+0.31 ± 0.08</b>   | (in base)             |
| Slack $S_{r(i)}$             | +0.12 ± 0.04          | −0.42 ± 0.31          |
| Centroid $d_{i,\text{cent}}$ | +0.10 ± 0.08          | <b>+0.48 ± 0.06</b>   |
| Density $\mu_{N/3}$          | +0.09 ± 0.12          | −0.08 ± 0.06          |
| Neighbor ranks               | +0.03 ± 0.05          | <b>+1.26 ± 0.09</b>   |
| Relative demand              | +0.03 ± 0.09          | −0.01 ± 0.02          |
| Density $\mu_5$              | +0.01 ± 0.01          | −0.02 ± 0.09          |
| Load share $\eta_i$          | +0.01 ± 0.04          | −0.42 ± 0.10          |
| Density $\mu_{N/10}$         | −0.03 ± 0.07          | −0.09 ± 0.09          |
| Load ratio $L_{r(i)}$        | −0.06 ± 0.08          | −0.44 ± 0.36          |

Table 2: Single-group signals, sorted by marginal value.  $\Delta_{\text{LOO}}$ : cost increase when the group is removed from the full set (positive = useful at the margin).  $\Delta_{\text{iso}}$ : cost reduction when the group alone is added to the static+meta base (positive = useful standalone). Bold: exceeds the noise floor. Static and meta belong to the isolation base and have no standalone entry.

appendix) are evaluated on top of the chosen set, individually and combined.

**Reference points.** The full 13-group representation reaches a mean final cost of  $16.678 \pm 0.026$ , against  $20.084 \pm 0.049$  for the static+meta base: the solution-dependent features cut cost by 3.41 routing units ( $\approx 17\%$ ). All configurations are trained with 3 seeds and evaluated seed-paired on the common test set; the pooled within-configuration seed std,  $\sigma \approx 0.18$ , serves as the significance floor for this sweep (it is inflated by the single high-variance no-static configuration; the typical seed std is  $\approx 0.05$ ).

**Single-group signals.** Table 2 and Figure 4 summarize the leave-one-out and isolation measurements (full plots in the appendix). Four groups clear the noise floor at the margin: static (whose removal is catastrophic, +3.84), detour, topology, and meta. In isolation, the geometric and relational signals carry almost all the standalone value — detour (+2.60), topology (+2.19), neighbor ranks (+1.26), centroid (+0.48)




Figure 4: Marginal (leave-one-out) against standalone (isolation) value per feature group; shaded bands mark the noise floor on each axis. The detour, topology and neighbor-rank groups are robustly useful by both criteria; the capacity signals (load ratio, slack, load share) sit in the bottom-left — weak at the margin and actively harmful in isolation.

| Set           | Groups | Cost   | Time (s) | $\Delta$ vs. full  |
|---------------|--------|--------|----------|--------------------|
| no load ratio | 12     | 16.648 | 173      | $-0.035 \pm 0.106$ |
| full set      | 13     | 16.683 | 173      | ref                |
| set7          | 7      | 17.020 | 129      | $+0.337 \pm 0.091$ |
| nodepct6      | 6      | 17.088 | 121      | $+0.405 \pm 0.077$ |
| centroid6     | 6      | 17.100 | 122      | $+0.417 \pm 0.049$ |
| core5         | 5      | 17.162 | 115      | $+0.479 \pm 0.075$ |

Table 3: Best-set validation at 5 seeds, seed-paired against the full set (noise floor  $\sigma \approx 0.057$ ; per-set seed std between 0.04 and 0.07). *core5* = static, meta, topology, neighbor ranks, detour; *set7* adds centroid and load share. Every coarsely pruned set is significantly worse; the only quality-neutral trim is dropping the load ratio alone.

— while the three capacity signals *hurt* when added alone to the base, and no density scale helps on its own. Thematic removals confirm the cluster picture (appendix table): the relational and geometry themes each cost  $\approx 0.6$  when removed, capacity  $+0.34$ , density only  $+0.15$ . Applied mechanically, the single-feature pruning rule would therefore retain only six groups (static, meta, topology, neighbor ranks, detour, centroid) and drop the density, capacity, and relative-demand groups.

**Joint validation overturns single-group pruning.** Because the leave-one-out and isolation protocols measure *single-group* signals, they are blind to collective value. The candidate sets were therefore re-trained head-to-head at 5 seeds (Table 3), on a batch whose noise floor is clean ( $\sigma \approx 0.057$ ). The verdict reverses the pruning rule: every down-pruned set is significantly worse than the full set, and cost degrades




Figure 5: Shared-encoder actor variant. Each node is encoded once into an embedding  $h_i$ ; two lightweight heads then score the two stages over the shared embeddings, either by concatenation or as a bilinear compatibility between the selected node and each candidate. An optional pooled summary  $g$  of all node embeddings (dashed) is appended to both heads’ inputs, giving every score access to instance-level context.

monotonically as groups are removed (core5  $\rightarrow$  set7  $\rightarrow$  full). The smallest penalty (set7,  $+0.337$ ) is already  $6\times$  the noise floor — the density and capacity groups that looked redundant one at a time contribute real joint value. The reversal is not about feature *count*: dropping the load ratio  $L_{r(i)}$  alone is a statistical tie with the full set ( $-0.035 \pm 0.106$  at 5 seeds, independently reproduced at 3 seeds with  $-0.059 \pm 0.076$ ), whereas additionally dropping the neighbor ranks costs an incremental  $+0.200 \pm 0.066$  — several times the floor. Pruning buys only compute: set7 trades 2% solution quality for a 25% speedup, an option under a tight inference budget but not the quality-optimal choice. We therefore retain the full feature set (the load ratio may be dropped at no cost; it brings no compute benefit either).

**Conditional pair descriptors.** On top of the retained set, the two second-stage conditioning signals were swept standalone and as a  $2 \times 2$  cross-product (3 seeds, noise floor  $\sigma \approx 0.071$ ). Both *hurt*: enabling the distance-rank descriptor costs  $+0.14 \pm 0.11$  and the insertion-cost descriptor  $+0.18 \pm 0.07$ , each adding  $\approx 20$  s of wall-clock; the  $2 \times 2$  main effects match the standalone deltas ( $+0.13$  and  $+0.17$ ) and the interaction is negligible ( $-0.02$ , within the floor), so the effects are additive and the both-on cell is the worst. The extra signal apparently distracts more than it informs: the per-node features already encode the local geometry the descriptors summarize. Both descriptors are disabled in all subsequent experiments.

## Policy and Critic Architecture

The architecture axis is swept in four steps, each inheriting the previous winner.

**Critic.** The baseline critic scores each node with a small MLP and averages the scalar outputs into the value estimate, which makes  $V_\phi$  a linear functional of the mean node representation. The alternative is the multi-statistic variant of the training section: nodes are encoded into embeddings, pooled with mean, max, and standard deviation, and mapped

| Variant                                                             | Cost          | $\Delta$ (paired)  | Time (s) |
|---------------------------------------------------------------------|---------------|--------------------|----------|
| <i>Critic</i> — reference: mean-pooled MLP, 16.693, 177 s           |               |                    |          |
| multi-statistic                                                     | <b>16.635</b> | −0.058             | 173      |
| <i>Actor capacity</i> — reference: $32 \times 1$ , 16.637, 173 s    |               |                    |          |
| largest cell ( $64 \times 2$ )                                      | 16.589        | $-0.049 \pm 0.039$ | 238      |
| <i>Actor structure</i> — reference: indep. scorers, $\approx 16.63$ |               |                    |          |
| + global context                                                    | 17.143        | $+0.508 \pm 0.496$ | 255      |
| shared enc. (best cell)                                             | 19.830        | $+3.196 \pm 0.103$ | 196      |
| <i>Distribution shaping</i> — reference: both off, 16.630, 175 s    |               |                    |          |
| logit clipping $C = 10$                                             | <b>16.457</b> | $-0.173 \pm 0.059$ | 177      |
| learnable temperature                                               | 16.580        | $-0.050 \pm 0.069$ | 176      |
| clipping + temp.                                                    | 16.481        | $-0.149 \pm 0.085$ | 179      |

Table 4: Architecture study summary (3 seeds per cell;  $\Delta$  = mean  $\pm$  std of the per-seed difference against the reference of its block, negative = better; bold = adopted changes). The critic  $\Delta$  is a difference of means because its reference arm carries an irregular seed set; the paired check on the two shared seeds gives  $-0.037 \pm 0.034$  (Table 11). Full grids, per-seed diagnostics, and plots are given in the appendix.

to the value by a nonlinear head — so dispersion information (clustered against scattered instances) survives the pooling.

**Capacity.** The embedding width and the number of hidden layers of the scoring networks are swept jointly, trading representational capacity against per-step inference cost.

**Structure.** The incumbent actor applies two independent MLPs directly to the raw feature vectors (Figure 2). We compare it with a *shared-encoder* variant (Figure 5): each node is encoded once into an embedding  $\mathbf{h}_i$ , and both stages are lightweight heads over the shared embeddings, the conditional score being computed either by a concatenation head over  $[\mathbf{h}_j \parallel \mathbf{h}_i]$  or as a bilinear compatibility  $\mathbf{q}^\top \mathbf{k} / \sqrt{d}$  between a query built from the selected node and a key per candidate. Orthogonally, a *global context* option pools the node representations into a summary  $g$  (mean, max, and standard deviation) appended to the input of both stages — the cheapest way to give the per-node scores access to instance-level information, at one pooling operation per step.

**Distribution shaping.** Finally, two flags tame the proposal distribution without adding capacity: *logit clipping*, which bounds the logits to  $C \tanh(z)$  before the softmax so the distribution cannot collapse to a near-deterministic choice (Kool, van Hoof, and Welling 2019), and a *learnable softmax temperature* per stage.

**Results.** Table 4 summarizes the four axes; every batch is seed-paired against its own incumbent and judged against its own pooled-seed-std noise floor. Two changes help, and neither adds capacity to the actor. The multi-statistic critic lowers cost by 0.058 routing units (floor 0.027) and is free at deployment, since the critic is discarded after training. Logit clipping is the only actor-side win: bounding the logits to  $10 \tanh(z)$  improves all three seeds, a seed-paired  $-0.173 \pm 0.059$  against a floor of 0.072, at unchanged compute. Everything that enlarges or rewires the actor is at best

neutral. The capacity grid is a plateau: the full spread from the smallest to the largest network is 0.079, no cell clears the floor (0.053) against the  $32 \times 1$  incumbent, and wall-clock rises by up to 38%; the faint monotone trend favors width over depth (main-effect spans 0.069 against 0.019). The pooled global context worsens *every* seed (paired  $+0.508 \pm 0.496$ , one seed degrading by +1.08) while adding 48% wall-clock — the instance-level summary distracts the per-node scores rather than informing them. The shared encoder is not a plateau but a cliff: all four cells of its grid land  $+3.20$  to  $+3.53$  above the reference, a 19% quality loss and more than  $18\times$  that batch’s noise floor (0.176), so forcing both selection stages through one trunk destroys the pair-scoring ability at this problem size; within the shared family the bilinear head beats concatenation ( $-0.31$ ) and the context toggle is neutral, but the best shared cell remains  $+3.20$  away. Finally, the learnable temperature is within noise on its own ( $-0.050 \pm 0.069$ ) and slightly degrades the clipped configuration when stacked (16.481 against 16.457, with doubled seed variance) — one more trainable knob for PPO to destabilize.

**Adopted configuration.** The retained policy keeps the incumbent actor on every structural axis — two independent per-stage scoring MLPs, 32-dimensional embeddings with a single hidden layer, no pooled context — and changes exactly two things: the multi-statistic critic and logit clipping at  $C = 10$  with a fixed temperature. The pattern is consistent: at  $N = 100$  the actor’s capacity and wiring are not the bottleneck; the leverage lies in the quality of the value estimate and in keeping the proposal distribution from collapsing — exactly the failure mode a near-deterministic policy induces inside a Metropolis loop, which needs proposal diversity to explore.

## Reward, Annealing, and Acceptance

We then validate the remaining components of the search loop.

**Reward shaping.** Let  $\Delta_t$  denote the cost improvement realized by the executed transition (zero when the move is rejected) and  $B_t$  the best cost found so far. The compared variants are: *immediate*,  $r_t = \Delta_t$  with a fixed penalty for infeasible proposals; *acceptance-aligned*, which rewards accepted improving moves by  $\Delta_t$  but deliberately assigns zero to accepted worsening moves (exploration is SA’s responsibility, not the policy’s fault) and a small constant penalty to rejected proposals; *best-so-far*,  $r_t = \max(0, B_{t-1} - B_t)$ , which only rewards new records; *terminal*, a single end-of-episode reward equal to the relative improvement over the initial cost; *mixtures*, either a fixed convex combination of the immediate and best-so-far signals or a scheduled interpolation that starts from the dense acceptance-aligned signal early in training and shifts toward the sparser best-so-far or terminal signals; and *step-penalty*, a negative per-step reward proportional to the current best cost, pressuring fast descent.

**Annealing.** We compare cooling schedule families (geometric decay  $T_{k+1} = \alpha T_k$  against step decay) and sweep the temperature range  $(T_0, T_{\text{final}})$ , which controls how much



| Variant                                                       | Cost        | $\Delta$ (paired)  | Time (s) |
|---------------------------------------------------------------|-------------|--------------------|----------|
| <i>Reward — reference: immediate, 16.459, 173 s</i>           |             |                    |          |
| fixed mixture                                                 | 16.460      | $+0.001 \pm 0.042$ | 176      |
| other dense variants                                          | 16.50–16.53 | $+0.04$ to $+0.07$ | 173–177  |
| step-penalty                                                  | 17.20–17.55 | $+0.74$ to $+1.09$ | 176      |
| terminal                                                      | 19.411      | $+2.952 \pm 0.983$ | 174      |
| <i>Cooling schedule — reference: geometric, 16.460, 179 s</i> |             |                    |          |
| step decay                                                    | 16.473      | $+0.013 \pm 0.034$ | 181      |
| <i>Metropolis — reference: on at train and test, 16.459</i>   |             |                    |          |
| off at test time                                              | 16.578      | $+0.119 \pm 0.025$ | —        |
| off at train time                                             | 16.432      | $-0.027 \pm 0.045$ | —        |
| off at both                                                   | 16.526      | $+0.068 \pm 0.050$ | —        |

Table 5: Reward, annealing, and acceptance summary (3 seeds per cell;  $\Delta$  = mean  $\pm$  std of the per-seed difference against the reference of its block, negative = better). All three axes confirm the incumbent configuration; the one effect that clears its noise floor unanimously is the cost of disabling the Metropolis acceptance at test time. Full rankings, per-seed diagnostics, and plots are given in the appendix.

worsening SA tolerates at the start and how strict it becomes at the end.

**Acceptance.** A  $2 \times 2$  ablation enables or disables the Metropolis criterion independently at training and at test time (its “off” arm accepts every proposed move), quantifying whether the learned proposal replaces the stochastic acceptance or complements it.

**Proposal mode at inference.** In standard RL applications, trained policies are often exploited greedily by selecting the most probable action. Our framework deviates from this practice: since the policy acts as the proposal distribution of an SA step, we retain stochastic sampling,  $a_t \sim \pi_\theta(\cdot | s_t)$ , at inference as well. A greedy policy would repeatedly propose the same “best” move, collapsing the diversity that the Metropolis criterion needs to balance exploration and exploitation; the ablation compares sampling against greedy arg-max exploitation.

**Initialization.** Each training episode starts from an initial solution produced by a constructive heuristic: random assignment, nearest neighbor, sweep, or single-customer routes. We compare training from each single heuristic against a progressively introduced mixture of all four — mixing exposes the policy to qualitatively different regions of the solution space and prevents it from overfitting to the basin of one construction heuristic. Every trained policy is additionally cross-tested from all initialization types, measuring how much the final quality depends on the starting point.

**Results.** Table 5 summarizes the three completed axes; as before, every comparison is seed-paired and judged against its batch’s pooled-seed-std noise floor. The reward sweep splits its nine candidates into two regimes. Every reward built on dense per-step improvement lands in a band only 0.07 wide: the incumbent immediate reward is best at 16.459, the fixed mixture is a statistical tie ( $+0.001 \pm 0.042$ , floor 0.042),




Figure 6: Curriculum ramp: the warm-up depth  $t_{\text{init}}$  grows from 0 to the warm-up strength  $\xi_{\text{CL}}$  over  $E$  training epochs, following a linear or sigmoid schedule, and stays saturated afterwards.

and the acceptance-aligned, best-so-far, and scheduled variants trail by at most 0.07 — the actor is largely insensitive to the exact shaping as long as the signal is dense. Removing the per-step credit is catastrophic: the terminal reward loses 2.95 routing units with seed swings of almost a full unit, and the step-penalty schemes lose 0.7 to 1.1 with similarly unstable training — the penalty fights the improvement signal instead of complementing it. The cooling-schedule comparison is a tie: the seed-paired gap between geometric and step decay is  $-0.013 \pm 0.034$  against a floor of 0.044, with the per-seed sign split, so the loop is insensitive to the shape of the cooling curve once its endpoints are fixed. The Metropolis cross gives the sharpest verdict: the criterion earns its keep at deployment, not at training. Removing the acceptance at test time costs  $+0.119 \pm 0.025$  over the incumbent-trained actor and  $+0.094 \pm 0.009$  over the greedy-trained one — several times the 0.034 floor and unanimous across seeds — whereas the training-time flag is within noise at the standard test-on deployment ( $+0.027 \pm 0.045$ , sign split). Only when deploying greedily does greedy training pay ( $+0.052 \pm 0.050$ ), a mild train/test-consistency effect outside the regime we deploy in.

**Confirmed configuration.** All three batches confirm the incumbent — immediate reward, geometric cooling, Metropolis acceptance kept during experience collection — so no configuration change is adopted, and they fix the deployment rule: the solver always runs with the Metropolis criterion on. The learned proposal complements the stochastic acceptance rather than replacing it; a greedy descent over the learned moves gets stuck in worse local optima, whatever trajectories the actor was trained on.

### Optimization Hyperparameters

The optimization knobs — actor and critic learning rates, entropy coefficient, and the rollout length (the number of SA steps collected per PPO update, which sets the horizon over which advantages are estimated) — are swept on the frozen configuration. The number of PPO epochs per collected batch is examined last and re-visited jointly with the curriculum in the final model, since the two interact.

### Curriculum Learning

Training episodes are much shorter than the search budgets used at deployment, so a policy trained only from raw con-



structive solutions rarely visits the fine-grained regime near local optima that dominates the end of a long search. The curriculum closes this gap by moving the *starting point*: before each collection phase, the batch of instances is first warmed up by the current policy for  $t_{\text{init}}$  steps, and  $t_{\text{init}}$  grows over training from 0 to a maximum warm-up strength  $\xi_{\text{CL}}$  over  $E$  epochs, following a linear or sigmoid ramp (Figure 6). Early in training the policy thus learns to repair poor solutions; later it increasingly learns to improve already-refined ones. A further variant draws a per-instance depth  $d_i \sim \mathcal{U}\{0, \dots, t_{\text{init}}\}$  instead of warming every instance to the same depth, so that a single batch covers the entire spectrum from raw initializations to deeply improved solutions.

The curriculum is validated with its own sequential protocol: (0) a multi-seed no-curriculum baseline, which also provides the noise floor used throughout this section; (1) a go/no-go comparison of a neutral mid-range curriculum against that baseline; (2) schedule shape (sigmoid against linear); (3) warm-up strength  $\xi_{\text{CL}}$ ; (4) ramp duration  $E$ ; (5) sigmoid steepness; (6) a mechanism ablation opposing the ramped warm-up to a constant warm-up of equal budget, separating the contribution of the progression itself; and (7) the per-instance random depth, hypothesized to pay off in robustness and out-of-distribution behavior more than in the in-distribution headline. The winning curriculum is confirmed against the baseline with 5 seeds.

## Final Model and Comparison with Baselines

All component winners, the best curriculum, and the tuned PPO budget are combined into the final model, trained with 5 seeds; the best checkpoint provides the headline numbers. We compare it against blind SA under matched conditions (same operator, budget, and schedule) and against the OR-Tools baseline, and we report the mean cost as a function of the SA step budget from  $10^2$  to  $10^6$  steps to characterize the anytime behavior of the guided search.

## Generalization and Scalability

Finally, we probe how far the node-centric representation carries across problem scales: models trained at each size are cross-tested at every other size, and the final model is evaluated on instances well beyond its training size with the same checkpoint, reporting both solution quality and runtime scaling.

## Conclusion

### Appendix: State Feature Definitions

This appendix defines every component of the node feature vector  $\mathbf{x}_i$  used by the policy and the critic. Throughout,  $\text{prev}(i)$  and  $\text{succ}(i)$  denote the nodes immediately before and after  $i$  in the current visiting order,  $r(i)$  the route currently serving  $i$ ,  $Q_{r(i)}$  the total load of that route, and  $d(\cdot, \cdot)$  the Euclidean distance. Figure 7 illustrates the geometric quantities. All features are normalized to comparable ranges, and each is computed locally at node  $i$ : the encoding contains no fixed-size global descriptor, which is what makes the representation applicable to any instance size.




Figure 7: Geometric features attached to a node  $i$  (orange) within its current route. The detour  $\Delta_i$  measures the extra length caused by visiting  $i$  between its neighbors;  $(\theta_i, d_{i0})$  locate  $i$  in polar coordinates around the depot;  $d_{i,\text{cent}}$  measures how far  $i$  lies from the center of gravity of its own route.

**Static instance features (6).** The 2D coordinates  $\mathbf{p}_i = (p_i^x, p_i^y) \in [0, 1]^2$ ; a depot indicator  $\mathbf{1}_{\{i=0\}}$ ; the polar angle  $\theta_i$  of the node around the depot; the distance to the depot  $d_{i0}$ , min-max normalized per instance; and the demand normalized by the vehicle capacity,  $q_i/Q$ . These features depend only on the instance, not on the current solution.

**Route topology (4).** The coordinates  $\mathbf{p}_{\text{prev}(i)}$  and  $\mathbf{p}_{\text{succ}(i)}$  of the node’s current neighbors in the visiting order. They give the policy direct access to the local shape of the route around each node.

**Neighbor distance-ranks (2).** For each of  $\text{prev}(i)$  and  $\text{succ}(i)$ , the rank of that neighbor in the list of all nodes sorted by distance from  $i$ , normalized to  $[0, 1]$  (the nearest node has rank  $\approx 0$ ). A well-placed node is linked to low-rank neighbors; a high rank signals a locally poor connection, independently of the absolute scale of the instance.

**Relative demand (1).** The demand normalized by the largest demand in the instance,  $q_i / \max_{j \in \mathcal{C}} q_j$ , complementing  $q_i/Q$  with a signal that is invariant to the capacity level.

**Spatial density (3).** The mean distance from  $i$  to its  $k$  nearest neighbors, evaluated at three scales:  $k = 5$ ,  $k = \lceil N/10 \rceil$ , and  $k = \lceil N/3 \rceil$ . These static statistics distinguish nodes in dense clusters from isolated ones.

**Local solution quality (2).** The *detour cost*

$$\Delta_i = d(\text{prev}(i), i) + d(i, \text{succ}(i)) - d(\text{prev}(i), \text{succ}(i)),$$

i.e., the extra length the tour pays for visiting  $i$  at its current position (equivalently, the saving obtained by removing it); and the distance  $d_{i,\text{cent}}$  from  $i$  to the centroid of the nodes of its route, which flags spatial outliers assigned to the wrong route.

**Capacity status (3).** The load ratio of the node’s route,  $L_{r(i)} = Q_{r(i)}/Q$ ; the remaining slack,  $S_{r(i)} = (Q - Q_{r(i)})/Q$ ; and the node’s share of its route load,  $\eta_i = q_i/Q_{r(i)}$ . These dynamic signals indicate whether a route is nearly saturated and how much a given node contributes to that saturation.

| Group           | Features                                                      | Dim |
|-----------------|---------------------------------------------------------------|-----|
| Static          | $\mathbf{p}_i, \mathbf{1}_{\{i=0\}}, \theta_i, d_{i0}, q_i/Q$ | 6   |
| Topology        | $\mathbf{p}_{\text{prev}(i)}, \mathbf{p}_{\text{succ}(i)}$    | 4   |
| Neighbor ranks  | rank of $\text{prev}(i), \text{succ}(i)$                      | 2   |
| Relative demand | $q_i / \max_j q_j$                                            | 1   |
| Spatial density | $\mu_5, \mu_{N/10}, \mu_{N/3}$                                | 3   |
| Local quality   | $\Delta_i, d_{i,\text{cent}}$                                 | 2   |
| Capacity status | $L_{r(i)}, S_{r(i)}, \eta_i$                                  | 3   |
| Meta            | temperature, remaining budget                                 | 2   |
| Total           |                                                               | 23  |

Table 6: Summary of the node feature vector  $\mathbf{x}_i$ .

| Operator  | Cost  | Feasibility | Cost  |
|-----------|-------|-------------|-------|
| Insertion | 20.00 | Masking     | 18.37 |
| 2-opt(*)  | 22.02 | Indirect    | 20.36 |
| Swap      | 26.03 | Rejection   | 29.31 |

Table 7: Main effects: mean final cost of each factor level averaged over the other factor. Insertion and masking win marginally as well as jointly.

**Search meta-features (2).** The current SA temperature, normalized to  $[0, 1]$  over the schedule range, and the fraction of the search budget remaining. They are identical across nodes of the same state and let the policy modulate its behavior along the annealing schedule (e.g., bolder restructuring at high temperature, fine repairs near the end).

**Conditional pair features (stage 2, optional).** In addition to the per-node vectors above, the second stage of the policy can receive descriptors of the candidate move itself, computed only for the selected first node  $u$  and every candidate partner  $v$ : the exact insertion cost  $d(v, u) + d(u, \text{succ}(v)) - d(v, \text{succ}(v))$  together with the net cost change of the full move (insertion cost minus the removal saving of  $u$ ), and the normalized distance-ranks of the edges the move would create. Because they are evaluated for a single  $u$ , these descriptors preserve the  $\mathcal{O}(N)$  per-step complexity.

## Appendix: Operator–Feasibility Study, Full Results

This appendix collects the complete diagnostics of the operator  $\times$  feasibility sweep summarized in Table 1. In the plot legends, the internal configuration names map to the paper terminology as `valid` = proactive masking, `free` = indirect representation (greedy split), `rm_depot` = reactive rejection.

### Appendix: Feature Ablation, Full Results

This appendix collects the complete plots and secondary tables of the feature-ablation study summarized in Tables 2 and 3. In the plot labels, the internal group names map to the paper terminology of Table 6 as `neighbor_rank` = neighbor distance-ranks, `demand_max` = relative demand, `density5/10%/33%` =  $\mu_5/\mu_{N/10}/\mu_{N/3}$ , `route_pct` =

| Operator  | Masking | Indirect | Rejection |
|-----------|---------|----------|-----------|
| Insertion | 173     | 172      | 152       |
| Swap      | 170     | 167      | 148       |
| 2-opt(*)  | 301     | 177      | 157       |

Table 8: Mean wall-clock time (s) to evaluate the full 10,000-instance test set for  $10^4$  SA steps. Rejection is cheapest because infeasible proposals are simply discarded; 2-opt(\*) under masking pays  $1.7\times$  the compute of every other cell for its suffix-exchange feasibility checks.




Figure 8: Mean final cost over the operator  $\times$  feasibility grid (green = better). The insertion/masking corner is the optimum; the rejection column is uniformly poor across all operators.

load ratio  $L_{r(i)}$ , slack = slack  $S_{r(i)}$ , and node\_pct = load share  $\eta_i$ .

## Appendix: Architecture Study, Full Results

This appendix collects the complete grids and per-seed diagnostics of the architecture study summarized in Table 4. In the plot labels, the internal names map to the paper terminology as `deepsets` = multi-statistic critic, `ff` = mean-pooled feed-forward critic, `seq` = the incumbent actor with independent per-stage scorers, `shared` = the shared-encoder variant of Figure 5, `gc` = pooled global context, `bil` = bilinear second-stage head (0 = concatenation), and networks are written width/depth, e.g. `64/2L`.

## Appendix: Reward, Annealing, and Acceptance, Full Results

This appendix collects the full rankings and per-seed diagnostics behind Table 5. All three batches follow the standard protocol (3 seeds, 10,000 SA steps at half precision on the Nazari  $N = 100$  test set, initial cost 57.87); the incumbent configuration reproduces at 16.459–16.460 across the three batches, matching the adopted cell of the architecture study (Table 14). In the plot labels, the internal reward names




Figure 9: Final cost grouped by operator, colored by feasibility strategy (error bars = seed std). Insertion wins under every strategy; swap is both worst and highest-variance.




Figure 10: Cost against evaluation time for the nine cells. Insertion/masking sits at the bottom left (best on both axes); 2-opt(\*)/masking spends  $1.7\times$  the compute and remains worse.

map to the paper terminology as `immediate` = immediate, `sa_aligned` = acceptance-aligned, `global_best` = best-so-far, `hybrid` = the fixed mixture ( $\alpha = 0.5$ ) of the immediate and best-so-far signals, `curriculum` and `curriculum_terminal` = the scheduled interpolations from acceptance-aligned toward best-so-far and terminal respectively, and `step_penalty_init` and `step_penalty_16` = the per-step penalty normalized by the initial cost and by a fixed constant close to the typical final cost.




Figure 11: Per-seed final cost for the four best cells. The winning cell is tightly clustered; seed variance grows quickly as quality drops.

| Theme removed                              | Cost   | $\Delta$ vs. full  |
|--------------------------------------------|--------|--------------------|
| Relational (topology + nbr. ranks)         | 17.277 | $+0.599 \pm 0.016$ |
| Geometry (detour + centroid)               | 17.274 | $+0.596 \pm 0.033$ |
| Capacity ( $L, S, \eta$ )                  | 17.021 | $+0.343 \pm 0.059$ |
| Density ( $\mu_5, \mu_{N/10}, \mu_{N/3}$ ) | 16.822 | $+0.145 \pm 0.033$ |

Table 9: Thematic ablation (3 seeds, seed-paired vs. the full set; positive  $\Delta$  = the theme helps). The relational and geometry clusters are the two most valuable; capacity clears the typical seed noise while density does not.

| Set                             | Groups | Cost   | Time (s) | $\Delta$ vs. full  |
|---------------------------------|--------|--------|----------|--------------------|
| full set                        | 13     | 16.678 | 179      | ref                |
| no load ratio                   | 12     | 16.619 | 166      | $-0.059 \pm 0.076$ |
| no load ratio,<br>no nbr. ranks | 11     | 16.819 | 154      | $+0.141 \pm 0.022$ |

Table 10: Targeted single-group refinements (3 seeds, seed-paired; mini-batch noise floor  $\approx 0.035$ ; per-set seed std between 0.02 and 0.05). Dropping the load ratio alone is a statistical tie with the full set; additionally dropping the neighbor ranks costs an incremental  $+0.200 \pm 0.066$ . The apparent 7% speedup of the 12-group set did not replicate at 5 seeds (Table 3) and is machine-load noise.




Figure 12: Leave-one-out ablation: cost increase when a single group is removed from the full set. Static dominates; detour, topology and meta are the next clearly useful groups. The shaded band is the noise floor.




Figure 13: Isolation ablation: gain from adding a single group to the static+meta base. Detour, topology and the neighbor ranks carry most of the standalone signal; the capacity trio harms the policy when added in isolation.




Figure 14: Thematic ablation (Table 9): removing the relational or geometry cluster costs  $\approx 0.6$  routing units each.

| Critic          | Cost               | Time (s) | Seeds     |
|-----------------|--------------------|----------|-----------|
| multi-statistic | $16.635 \pm 0.038$ | 173      | {0, 1, 2} |
| feed-forward    | $16.693 \pm 0.006$ | 177      | {1, 2, 2} |

Table 11: Critic comparison (3 runs per arm; noise floor  $\approx 0.027$ ). The feed-forward arm is missing seed 0 and duplicates seed 2, so the main-text  $\Delta$  is a difference of means; restricted to the two shared seeds, the paired delta is  $-0.037 \pm 0.034$  — the same direction, suggestive on its own but consistent with the unpaired gap.

| Net           | Cost               | Time (s) | $\Delta$ vs. $32 \times 1$ |
|---------------|--------------------|----------|----------------------------|
| $64 \times 2$ | $16.589 \pm 0.004$ | 238      | $-0.049 \pm 0.039$         |
| $64 \times 1$ | $16.591 \pm 0.084$ | 207      | $-0.047 \pm 0.060$         |
| $32 \times 2$ | $16.601 \pm 0.008$ | 187      | $-0.037 \pm 0.030$         |
| $32 \times 1$ | $16.637 \pm 0.036$ | 173      | ref                        |
| $16 \times 2$ | $16.649 \pm 0.083$ | 165      | $+0.011 \pm 0.093$         |
| $16 \times 1$ | $16.668 \pm 0.042$ | 160      | $+0.031 \pm 0.063$         |

Table 12: Actor capacity grid (embedding width  $\times$  hidden layers), sorted by mean cost;  $\Delta$  is seed-paired against the  $32 \times 1$  incumbent. No cell clears the noise floor ( $\approx 0.053$ ); the width main effect spans 0.069 against 0.019 for depth.

| Stage-2 head | Context | Cost   | Time (s) | $\Delta$ vs. ref   |
|--------------|---------|--------|----------|--------------------|
| bilinear     | on      | 19.830 | 196      | $+3.196 \pm 0.103$ |
| bilinear     | off     | 19.860 | 168      | $+3.226 \pm 0.163$ |
| concat       | off     | 20.154 | 192      | $+3.520 \pm 0.251$ |
| concat       | on      | 20.160 | 236      | $+3.526 \pm 0.099$ |

Table 13: Shared-encoder grid, sorted by mean cost;  $\Delta$  is seed-paired against the incumbent actor re-run with context off ( $16.633 \pm 0.036$ , 180 s). Per-cell seed std between 0.06 and 0.26; noise floor  $\approx 0.176$ . Within the shared family the bilinear head beats concatenation (main effect  $-0.31$ ) and the context toggle is neutral ( $-0.01$ ), but every cell is  $\approx 3$  units worse than the reference.

| Clipping | Temp.   | Cost          | Time (s) | $\Delta$ vs. ref   |
|----------|---------|---------------|----------|--------------------|
| $C = 10$ | fixed   | <b>16.457</b> | 177      | $-0.173 \pm 0.059$ |
| $C = 10$ | learned | 16.481        | 179      | $-0.149 \pm 0.085$ |
| off      | learned | 16.580        | 176      | $-0.050 \pm 0.069$ |
| off      | fixed   | 16.630        | 175      | ref                |

Table 14: Distribution-shaping grid, sorted by mean cost;  $\Delta$  is seed-paired against the both-off incumbent. Per-cell seed std between 0.04 and 0.11; noise floor  $\approx 0.072$ . Only logit clipping clears the floor; the learnable temperature does not help it (clipping main effect  $-0.136$ , temperature  $-0.014$ ).

| Reward                           | Cost                                 | $\Delta$ (paired)  |
|----------------------------------|--------------------------------------|--------------------|
| immediate (incumbent)            | <b>16.459 <math>\pm</math> 0.039</b> | ref                |
| fixed mixture                    | 16.460 $\pm$ 0.017                   | +0.001 $\pm$ 0.042 |
| sched. $\rightarrow$ best-so-far | 16.503 $\pm$ 0.078                   | +0.043 $\pm$ 0.053 |
| best-so-far                      | 16.510 $\pm$ 0.014                   | +0.051 $\pm$ 0.049 |
| acceptance-aligned               | 16.516 $\pm$ 0.010                   | +0.057 $\pm$ 0.035 |
| sched. $\rightarrow$ terminal    | 16.531 $\pm$ 0.050                   | +0.072 $\pm$ 0.023 |
| step-penalty (const.)            | 17.198 $\pm$ 0.386                   | +0.739 $\pm$ 0.425 |
| step-penalty (init)              | 17.549 $\pm$ 0.716                   | +1.090 $\pm$ 0.707 |
| terminal                         | 19.411 $\pm$ 0.955                   | +2.952 $\pm$ 0.983 |

Table 15: Full reward ranking, sorted by mean cost;  $\Delta$  is seed-paired against the incumbent. The horizontal rule separates the stable cluster (six rewards within 0.15 of the incumbent) from the broken sparse and penalty-based rewards. Wall-clock is flat across all arms (173–177 s). The noise floor is the pooled seed std over the stable cluster,  $\approx 0.042$ ; the global pooled std (0.420) is inflated by the unstable training of the broken arms and is not used as a threshold.

| Schedule              | Cost               | Time (s) |
|-----------------------|--------------------|----------|
| geometric (incumbent) | 16.460 $\pm$ 0.040 | 179      |
| step decay            | 16.473 $\pm$ 0.048 | 181      |

Table 16: Cooling-schedule comparison. Both schedules share the same temperature range and differ only in the shape of the decay: geometric,  $T_{k+1} = \alpha T_k$ , against a piecewise-constant schedule dropping  $T$  at one half, two thirds, and three quarters of the horizon. The seed-paired  $\Delta$  (geometric – step) is  $-0.013 \pm 0.034$  against a noise floor of 0.044, with one of three seeds favoring step decay — a statistical tie, and wall-clock is identical within jitter.

|                                                                             | Test on (default)  | Test off           |
|-----------------------------------------------------------------------------|--------------------|--------------------|
| train on (incumbent)                                                        | 16.459 $\pm$ 0.040 | 16.578 $\pm$ 0.050 |
| train off (greedy)                                                          | 16.432 $\pm$ 0.010 | 16.526 $\pm$ 0.020 |
| <i>test-time effect (off – on):</i> +0.119 $\pm$ 0.025 / +0.094 $\pm$ 0.009 |                    |                    |
| <i>training effect (on – off):</i> +0.027 $\pm$ 0.045 / +0.052 $\pm$ 0.050  |                    |                    |

Table 17: Metropolis  $2 \times 2$  cross (mean  $\pm$  seed std). Effects are seed-paired: the test-time effect is reported at train-on / train-off and the training effect at test-on / test-off. Against the 0.034 noise floor, both test-time effects are unanimous across seeds and several times the floor; the training effect at the standard test-on deployment is a tie (sign split across seeds), and only under greedy deployment does greedy training clear the floor — a mild train/test-consistency effect. The greedy-trained arms also show markedly lower seed variance (0.010/0.020 against 0.040/0.050).




Figure 15: All 30 ablation configurations ranked by mean final cost. Leave-one-out configurations (dark blue) cluster tightly near the full-set reference (red dashed); isolation configurations (light blue) sit near the static+meta base (dotted).






Figure 16: Per-seed final cost for the key reference configurations (red bars = means). Seed variance is tiny except for the no-static configuration, which inflates the pooled noise floor of the 3-seed sweep.




Figure 17: Best-set validation (Table 3): candidate sets ranked at 5 seeds. The 12-group set (gold) and the full set (green) tie at the top; every coarsely pruned set (blue) is clearly worse.




Figure 18: Best-set validation, per-seed view. Left: per-seed final cost — the 12-group and full-set clusters overlap and sit cleanly below every pruned set. Right: seed-paired  $\Delta$  vs. the full set — the pruned sets (red) clear the noise floor by a wide margin; the 12-group set (green) sits below zero.




Figure 19: Best-set validation, cost-compute view. Left: pruning moves down-left (cheaper but worse); the 12-group set matches the full set on both axes. Right: cost against number of feature groups — the pruned family degrades monotonically as groups are removed.




Figure 20: Targeted refinements (Table 10). Left: per-seed final cost — dropping the load ratio overlaps the full set, while additionally dropping the neighbor ranks lifts the whole cluster. Right: seed-paired  $\Delta$  vs. the full set — the load-ratio removal (green) sits inside the noise floor, the neighbor-rank removal (red) clears it decisively.




Figure 21: Conditional pair descriptors. Left:  $2 \times 2$  grid of mean final cost — the both-off cell (plain retained set) is best, both-on is worst. Right: standalone deltas and  $2 \times 2$  main effects — both descriptors sit clearly beyond the noise floor on the harmful side; the interaction is within it.




Figure 22: Critic comparison (Table 11). Left: per-seed final cost (red bars = means); the multi-statistic cluster sits below the feed-forward one. Right: mean cost against evaluation time — the multi-statistic critic concedes nothing on either axis.




Figure 23: Actor capacity grid (Table 12). Left: mean final cost and wall-clock per cell (green = better). Right: seed-paired  $\Delta$  against the  $32 \times 1$  incumbent — every cell falls inside the shaded noise band, so no capacity change is a significant win.




Figure 24: Cost against compute across the six capacity cells, colored by embedding width. Quality improves only marginally along a steeply rising compute axis — the hallmark of a capacity plateau.




Figure 25: Global-context toggle. Left: per-seed final cost (black bars = means); the context-on arm is both worse and far more variable. Right: per-seed  $\Delta$  from enabling context — every seed is positive (worse), one dramatically so, so the degradation is not a variance artefact.




Figure 26: Shared encoder against independent scorers (Table 13). Left: per-seed final cost, incumbent reference (green) against the four shared cells (red); the dashed line marks the reference mean, and every shared seed sits far above it. Right: the shared  $2 \times 2$  grid — the bilinear head beats concatenation and context is roughly neutral, but all four cells are  $\approx 3$  units worse than the reference.




Figure 27: Distribution shaping (Table 14). Left:  $2 \times 2$  grid of mean final cost and time (green = better); the clipping-on, fixed-temperature cell is the best. Right: seed-paired  $\Delta$  against the both-off incumbent — only logit clipping escapes the noise band, and stacking the learnable temperature on top of it does not help.




Figure 28: Full reward ranking (Table 15; mean  $\pm$  seed std, best at top). Gold = incumbent, green = the stable cluster within 0.15 of it, red = degraded. The three red bars (terminal and the two step-penalty schemes) are a different regime entirely: they fail to train a useful actor and carry runaway seed variance.






Figure 29: Zoom on the stable reward cluster. Left: per-seed final cost (black bars = means); the arms overlap heavily. Right: seed-paired  $\Delta$  against the incumbent immediate reward — the fixed mixture sits inside the noise band (a tie), the remaining dense variants are marginally but consistently above zero.




Figure 30: Cooling-schedule comparison (Table 16). Left: per-seed final cost for the two schedules (black bars = means); the clusters overlap heavily. Right: per-seed paired  $\Delta$  (geometric – step) — every bar lies inside the noise band and the sign is not consistent, the signature of a tie rather than an effect.




Figure 31: Metropolis train  $\times$  test cross (Table 17). Left:  $2 \times 2$  heatmap of mean final cost (greener = better). Center: cost of removing test-time Metropolis (off – on) at each training setting — both bars are well past the noise band, the dominant effect. Right: training-flag effect (on – off) at each deployment — the test-on bar sits inside the band (a tie), test-off just clears it.

## References

- Correia, A. H. C.; Worrall, D. E.; and Bondesan, R. 2023. Neural Simulated Annealing. In *Proceedings of The 26th International Conference on Artificial Intelligence and Statistics*, 4946–4962. PMLR. ISSN: 2640-3498.
- Furnon, V.; and Perron, L. 2024. OR-Tools Routing Library.
- Kirkpatrick, S.; Gelatt, C. D.; and Vecchi, M. P. 1983. Optimization by Simulated Annealing. 220.
- Kool, W.; van Hoof, H.; and Welling, M. 2019. Attention, Learn to Solve Routing Problems! In *International Conference on Learning Representations*.
- Ma, Y.; Li, J.; Cao, Z.; Song, W.; Zhang, L.; Chen, Z.; and Tang, J. 2021. Learning to Iteratively Solve Routing Problems with Dual-Aspect Collaborative Transformer. In *Advances in Neural Information Processing Systems*, volume 34, 11096–11107. Curran Associates, Inc.
- Nazari, M.; Oroojlooy, A.; Snyder, L.; and Takac, M. 2018. Reinforcement Learning for Solving the Vehicle Routing Problem. In *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc.
- Schulman, J.; Wolski, F.; Dhariwal, P.; Radford, A.; and Klimov, O. 2017. Proximal Policy Optimization Algorithms. ArXiv:1707.06347 [cs].
- Vidal, T. 2016. Technical note: Split algorithm in  $O(n)$  for the capacitated vehicle routing problem. *Computers & Operations Research*, 69: 40–47.